

MILAN AI WEEK • 2026

RED TEAM

AGENT

Adversarial AI that attacks your strategic documents before the market does.

The Problem

Every major decision goes unchallenged.



The document lands

A pitch deck, an IPO prospectus, an acquisition memo. It looks solid. It has been iterated internally for weeks.



The room agrees

Everyone in the room has an incentive to say yes. The banker wants the deal. The founder believes the story. The board trusts the team.



The market disagrees

Six months later, the assumptions were wrong. The risks were visible. Nobody asked the hard questions.

\$47B → \$44.5M

WeWork valuation: January 2019 to November 2023

A 99.9% loss of value. The risks were in the S-1.

The WeWork S-1 contained every warning sign.

CRITICAL

– \$17.9B in obligations vs. \$4B in revenue commitments. Members exit on 30 days' notice.

HIGH

– same model, profitable, 10x the locations, valued 12.7x lower. Nobody asked why.

CRITICAL

– "Community Adjusted EBITDA" turned a \$933M loss into a \$233M apparent profit.

HIGH

– members on monthly contracts. 15-year leases. Asymmetric fragility.

CRITICAL

– Neumann leased buildings he personally owned back to his own company.

MEDIUM

– no permanent CFO for 12 months before the IPO attempt.

All of this was public information. No one ran a structured adversarial analysis.

The Solution

Red Team Agent

Five specialized AI agents attack your document simultaneously – each from a different adversarial perspective.

They identify financial, legal, market, competitive, and operational vulnerabilities. A synthesis agent scores the overall risk and surfaces the questions the board must answer before proceeding.

Risk Score 0-100

Structured findings

Remediation roadmap



CFO Agent

attacking

Financial assumptions & metric integrity



Market Agent

attacking

Demand, moat, and customer fragility



Legal Agent

3 found

Governance & conflicts of interest




Competitor Agent

2 found

Competitive landscape & barriers to entry

How It Works



PDF, Markdown, or load a
demo

Document + company context
chunked and tagged

Each attacks from a different
angle simultaneously

Patterns, cross-agent risks, and
verdict

Risk score, findings,
remediation roadmap

End-to-end in under

The 5 Adversarial Agents

Each agent has a single job: find the flaw.



CFO

Attacks financial assumptions, metric integrity, burn rate, and valuation logic.

FINANCIAL ADVERSARY



Market

Challenges demand projections, switching costs, and identifies where the moat is quicksand.

DEMAND SKEPTIC



Legal

Finds conflicts of interest and governance structures that protect insiders at everyone else's expense.

**GOVERNANCE
WATCHDOG**



Competitor

Explains exactly how and why the competition wins, and why barriers to entry are lower than stated.

RIVAL CEO



Execution

Identifies why this specific organization will never pull this specific plan off.

COO SKEPTIC

All 5 agents run

. First result in ~3 seconds.

RAG-Powered Company Context

Your documents make the analysis yours.

Red Team Agent retrieves context from your internal knowledge base before each agent runs. The findings are grounded in your company's actual data – not generic assumptions.



Strategy docs

Internal memos, board decks, past analyses



Financial data

Historical P&L, budgets, forecasts

HOW RAG WORKS

1

Document is chunked and tagged by domain (financial, legal, market, competitive, operational)

2

Each agent retrieves only the chunks relevant to its domain before generating findings

3

Internal context (your data) is layered on top of the decision document for cross-referenced analysis

4

Every finding is grounded in cited source passages – not hallucinated

Integrations

Connect your existing knowledge base.



Notion

Wikis, docs, project notes



Google Drive

Files, spreadsheets, slides



Salesforce

CRM data, deal history



Confluence

Engineering, product
docs



Slack

Team conversations,
decisions



SharePoint

Corporate docs, policies



● Connected ● Coming soon · Or upload files directly – PDF, Markdown, TXT

The Output

A complete adversarial report in 90 seconds.

RISK SCORE & VERDICT

94 /100

DO NOT PROCEED

Overall risk score derived from agent findings, weighted by severity and cross-agent patterns.

STRUCTURED FINDINGS

CRITICAL Lease gap: \$13.9B exposed

CRITICAL Invented financial metric

HIGH Zero switching cost

HIGH IWG: same model, 12.7x cheaper

+ cited sources for each finding

3 QUESTIONS TO ANSWER

- 1.** *What is the occupancy breakeven in a 20% churn recession?*
- 2.** *How many Neumann-WeWork transactions weren't disclosed?*
- 3.** *Why does IWG trade at 12.7x less?
Answer must be quantitative.*

"WeWork burned \$47 billion in valuation in four years."

The Red Team Agent would have flagged it in 90 seconds.

MILAN AI WEEK • 2026

Ask the hard questions first.

Red Team Agent – adversarial AI analysis for strategic documents.

Try the demo ↗

github.com/galatro/hackaton-ai-week-2026 ↗